

# PredictNextMinLoad

October 26, 2017

```
In [3]: from pyspark.sql import SparkSession
```

```
# Build the SparkSession
spark = SparkSession.builder \
    .master("local") \
    .appName("Machine learning for Load Prediction") \
    .config("spark.executor.memory", "1gb") \
    .getOrCreate()
sc = spark.sparkContext
# Load the data by creating rdd
rdd = sc.textFile('/home/hassan/Side_Projects/WeblogChallenge/data/2015_07_22_mktplace_s
# split the data into columns
rdd = rdd.map(lambda line: line.split(" "))
# =====
# Manipulating data
# =====
from pyspark.sql import Row
from pyspark.sql.types import *
from pyspark.sql.functions import *

#Map the RDD to a DF for better performance
mainDF = rdd.map(lambda line: Row(timestamp=line[0], ipaddress=line[2].split(':')[0],url
# convert timestamps from string to timestamp datatype
mainDF = mainDF.withColumn('timestamp', mainDF['timestamp'].cast(TimestampType()))

#get count of hit within window of 60 every second
loadperMinDF = mainDF.select(window("timestamp", "60 seconds").alias('timewindow'),'time
# get count of hit per IP
countdDF = mainDF.select(window("timestamp", "60 seconds").alias('timewindow'),'timestamp')
countdDF.show(20)
```

```
+-----+-----+-----+
|          timewindow|          ipaddress|HitperMin|
+-----+-----+-----+
|[2015-07-22 05:00...|117.200.191.192|          1|
|[2015-07-22 05:00...| 106.77.203.224|          1|
|[2015-07-22 05:00...| 122.15.164.218|          3|
|[2015-07-22 05:00...| 107.150.4.153|          1|
```

```
| [2015-07-22 05:00... | 103.242.156.9 | 1 |
| [2015-07-22 05:00... | 106.67.99.24 | 1 |
| [2015-07-22 05:00... | 107.167.108.212 | 2 |
| [2015-07-22 05:00... | 122.166.231.76 | 1 |
| [2015-07-22 05:00... | 182.68.216.254 | 2 |
| [2015-07-22 05:00... | 117.234.213.177 | 1 |
| [2015-07-22 05:00... | 49.205.99.169 | 1 |
| [2015-07-22 05:00... | 117.211.43.234 | 2 |
| [2015-07-22 05:00... | 49.207.236.65 | 1 |
| [2015-07-22 05:00... | 107.167.108.131 | 2 |
| [2015-07-22 05:00... | 128.185.3.222 | 1 |
| [2015-07-22 05:01... | 115.113.117.48 | 14 |
| [2015-07-22 05:01... | 1.39.13.51 | 3 |
| [2015-07-22 05:01... | 203.122.41.18 | 5 |
| [2015-07-22 05:01... | 117.196.181.12 | 6 |
| [2015-07-22 05:01... | 116.203.5.226 | 5 |
```

```
+-----+-----+-----+
```

only showing top 20 rows

```
In [4]: # computing mean ,std and max of hit counts per IP within 60 second window as
# features of each 60 second window
```

```
# these features can be used for perdicting the next load in next minute
```

```
Feature1 = countdDF.groupBy("timewindow").agg(stddev('HitperMin').alias("stdOfHitPerMinPerIP"))
```

```
Feature2 = countdDF.groupBy("timewindow").agg(mean('HitperMin').alias("meanOfHitPerMinPerIP"))
```

```
Feature3 = countdDF.groupBy("timewindow").agg(max('HitperMin').alias("maxOfHitPerMinPerIP"))
```

```
Features = Feature1.join(Feature2,["timewindow"])
```

```
Features = Features.join(Feature3,["timewindow"])
```

```
Features = Features.join(loadperMinDF,["timewindow"])
```

```
Features = Features.orderBy('timewindow', ascending=True)
```

```
Features.show(20,False)
```

```
+-----+-----+-----+
|timewindow|stdOfHitPerMinPerIP|meanOfHitPerMinPerIP|maxOfHitPerMinPerIP|
+-----+-----+-----+
| [2015-07-21 22:40:00.0,2015-07-21 22:41:00.0] |37.63654992233117 |6.183619550858652 |978
| [2015-07-21 22:41:00.0,2015-07-21 22:42:00.0] |45.16269677704943 |8.580278128950695 |846
| [2015-07-21 22:42:00.0,2015-07-21 22:43:00.0] |41.00957913182911 |7.097839898348157 |915
| [2015-07-21 22:43:00.0,2015-07-21 22:44:00.0] |21.876291841819434 |5.738181818181818 |366
| [2015-07-21 22:44:00.0,2015-07-21 22:45:00.0] |19.665877050690277 |5.591397849462366 |374
| [2015-07-21 22:45:00.0,2015-07-21 22:46:00.0] |2.0763215336529917 |1.8888888888888888 |18
| [2015-07-22 01:09:00.0,2015-07-22 01:10:00.0] |NaN |1.0 |1
| [2015-07-22 01:10:00.0,2015-07-22 01:11:00.0] |13.95908526268474 |5.097930338213024 |355
| [2015-07-22 01:11:00.0,2015-07-22 01:12:00.0] |16.928535755315508 |6.049766718506999 |346
| [2015-07-22 01:12:00.0,2015-07-22 01:13:00.0] |22.19877473979011 |6.37948984903696 |539
| [2015-07-22 01:13:00.0,2015-07-22 01:14:00.0] |25.967580688670626 |6.771067415730337 |467
```

[2015-07-22 01:14:00.0,2015-07-22 01:15:00.0]	18.09106237421149	5.919556840077071	575
[2015-07-22 01:15:00.0,2015-07-22 01:16:00.0]	6.7884838863006856	2.982490272373541	75
[2015-07-22 02:54:00.0,2015-07-22 02:55:00.0]	0.0	1.0	1
[2015-07-22 02:55:00.0,2015-07-22 02:56:00.0]	24.397729320063686	5.4246753246753245	1058
[2015-07-22 02:56:00.0,2015-07-22 02:57:00.0]	13.954862803139653	5.335841584158416	493
[2015-07-22 02:57:00.0,2015-07-22 02:58:00.0]	11.228685780753388	5.507252382925818	152
[2015-07-22 02:58:00.0,2015-07-22 02:59:00.0]	10.784985916123745	5.006719865602688	274
[2015-07-22 02:59:00.0,2015-07-22 03:00:00.0]	11.13514600581905	5.062182741116751	278
[2015-07-22 03:00:00.0,2015-07-22 03:01:00.0]	4.046688058258602	2.400749063670412	59

only showing top 20 rows

In [5]: # Divide hit counts by 60 to become hit per second

```
Features = Features.withColumn("stdOfHitPerSecPerIP", col("stdOfHitPerMinPerIP")/60.0) \
    .withColumn("meanOfHitPerSecPerIP", col("meanOfHitPerMinPerIP")/60.0) \
    .withColumn("maxOfHitPerSecPerIP", col("maxOfHitPerMinPerIP")/60.0) \
    .withColumn("HitperSec", col("HitperMin")/60.0)
Features = Features.select("timewindow", "stdOfHitPerSecPerIP", "meanOfHitPerSecPerIP", "maxOfHitPerSecPerIP", "HitperSec")
Features.show(5)
```

timewindow	stdOfHitPerSecPerIP	meanOfHitPerSecPerIP	maxOfHitPerSecPerIP	HitperSec
[2015-07-21 22:40...]	0.6272758320388528	0.1030603258476442	16.3	78.01666666
[2015-07-21 22:41...]	0.7527116129508239	0.14300463548251158	14.1	113.11666666
[2015-07-21 22:42...]	0.6834929855304852	0.11829733163913596	15.25	
[2015-07-21 22:43...]	0.3646048640303239	0.09563636363636364	6.1	
[2015-07-21 22:44...]	0.3277646175115046	0.0931899641577061	6.233333333333333	

only showing top 5 rows

In [6]: # get id for each window

```
Features = Features.withColumn("tagId", monotonically_increasing_id().cast("double"))
Features.show(10)
```

timewindow	stdOfHitPerSecPerIP	meanOfHitPerSecPerIP	maxOfHitPerSecPerIP	HitperSec	tagId
[2015-07-21 22:40...]	0.6272758320388528	0.1030603258476442	16.3	78.01666666	0.016666666666666666
[2015-07-21 22:41...]	0.7527116129508239	0.14300463548251158	14.1	113.11666666	0.03333333333333333
[2015-07-21 22:42...]	0.6834929855304852	0.11829733163913596	15.25		0.05
[2015-07-21 22:43...]	0.3646048640303239	0.09563636363636364	6.1		0.06666666666666666
[2015-07-21 22:44...]	0.3277646175115046	0.0931899641577061	6.233333333333333		0.08333333333333333
[2015-07-21 22:45...]	0.03460535889421653	0.03148148148148148	0.3	5.38333333	0.1
[2015-07-22 01:09...]	NaN	0.016666666666666666	0.016666666666666666	0.016666666666666666	0.11666666666666666

```
| [2015-07-22 01:10... | 0.23265142104474565 | 0.08496550563688372 | 5.916666666666667 | 168.31666666666666 |
| [2015-07-22 01:11... | 0.2821422625885918 | 0.10082944530844998 | 5.766666666666667 | 168.31666666666666 |
| [2015-07-22 01:12... | 0.36997957899650186 | 0.10632483081728267 | 8.983333333333333 | 168.31666666666666 |
+-----+-----+-----+-----+-----+
only showing top 10 rows
```

stdOfHitPerSecPerIP	meanOfHitPerSecPerIP	maxOfHitPerSecPerIP	HitperSec	LoadInNextSec
0.6272758320388528	0.1030603258476442	16.3	78.01666666666667	113.11666666666667
0.7527116129508239	0.14300463548251158	14.1	113.11666666666666	
0.6834929855304852	0.11829733163913596	15.25	93.1	
0.3646048640303239	0.09563636363636364	6.1	78.9	
0.3277646175115046	0.0931899641577061	6.233333333333333	78.0	5.383333333333333
0.03460535889421653	0.03148148148148148	0.3	5.383333333333333	0.016666666666666666
NaN	0.016666666666666666	0.016666666666666666	0.016666666666666666	168.31666666666666
0.23265142104474565	0.08496550563688372	5.916666666666667	168.31666666666666	
0.2821422625885918	0.10082944530844998	5.766666666666667	194.5	
0.36997957899650186	0.10632483081728267	8.983333333333333	204.25	

only showing top 10 rows

```

from pyspark.ml.regression import LinearRegression, RandomForestRegressor

rf = RandomForestRegressor(numTrees=100, maxDepth=10)
linearModel = rf.fit(train_data)

#lr = LinearRegression(labelCol="label", maxIter=100, regParam=0.3, elasticNetParam=0.8)
# Fit the data to the model
#linearModel = lr.fit(train_data)

predicted = linearModel.transform(test_data)
predictions = predicted.select("prediction").rdd.map(lambda x: x[0])
labels = predicted.select("label").rdd.map(lambda x: x[0])
predictionAndLabel = predictions.zip(labels).collect()

```

```

In [23]: import numpy as np
         error = []
         for a in predictionAndLabel:
             error.append(np.abs(a[0]-a[1]))

         print 'mean abs error is: ', np.mean(error)

```

mean abs error is: 65.400781239

```

In [28]: # predicting the load for the next minute
         # because the last data record is the last 60 second of data
         # we predict for the last record for predicting the next minute load

         with_id = NextloadDf.withColumn("_id", monotonically_increasing_id())
         i = with_id.select(max("_id")).first()[0]
         last_item = with_id.where(col("_id") == i).drop("_id")
         input_data = last_item.rdd.map(lambda x: (x[4], DenseVector(x[:4])))
         dataframeInputdata = spark.createDataFrame(input_data, ["label", "features"])
         predicted = linearModel.transform(dataframeInputdata)
         predictions = predicted.select("prediction").rdd.map(lambda x: x[0])
         labels = predicted.select("label").rdd.map(lambda x: x[0])
         predictions = predictions.collect()
         predictions[:]

```

Out[28]: [69.80983333333323]

In [ ]: