

# PredictSessionLengthURLVisits

October 26, 2017

```
In [3]: from pyspark.sql import SparkSession

# Build the SparkSession
spark = SparkSession.builder \
    .master("local") \
    .appName("Predict Session Length for given IP") \
    .config("spark.executor.memory", "1gb") \
    .getOrCreate()
sc = spark.sparkContext

# Load the data by creating rdd
rdd = sc.textFile('/home/hassan/Side_Projects/WeblogChallenge/data/2015_07_22_mktplace_s
# split the data into columns
rdd = rdd.map(lambda line: line.split(" "))

# =====
# Manipulating data
# =====
from pyspark.sql import Row
from pyspark.sql.types import *
from pyspark.sql.functions import *

#Map the RDD to a DF for better performance
mainDF = rdd.map(lambda line: Row(timestamp=line[0], ipaddress=line[2].split(':')[0],url=
# convert timestamps from string to timestamp datatype
mainDF = mainDF.withColumn('timestamp', mainDF['timestamp'].cast(TimestampType()))
# sessionizing data based on 15 min fixed window time
# assign an Id to each session
SessionDF = mainDF.select(window("timestamp", "15 minutes").alias('FixedTimeWindow'),'ti
SessionDF = SessionDF.withColumn("SessionId", monotonically_increasing_id())
# join the time stamps and url to the Sessionized DF
dfWithTimeStamps = mainDF.select(window("timestamp", "15 minutes").alias('FixedTimeWindo
SessionDF = dfWithTimeStamps.join(SessionDF,['FixedTimeWindow','ipaddress'])
# Finding the first hit time of each ip for each session and join in to our session df
FirstHitTimeStamps = SessionDF.groupBy("SessionId").agg(min("timestamp").alias('FristHit
SessionDF = FirstHitTimeStamps.join(SessionDF,['SessionId'])
timeDiff = (unix_timestamp(SessionDF.timestamp)-unix_timestamp(SessionDF.FristHitTime))
```

```

SessionDF = SessionDF.withColumn("timeDiffwithFirstHit", timeDiff)
tmpdf = SessionDF.groupBy("SessionId").agg(max("timeDiffwithFirstHit").alias("SessionDuration"))
SessionDF = SessionDF.join(tmpdf,['SessionId'])

```

```

# for any given IP if we don't have any previous log from that IP
# the prediction for it's session length is the average session length
meandf = SessionDF.groupBy().avg('SessionDuration')
meandf.show()

```

```

+-----+
|avg(SessionDuration)|
+-----+
| 141.58578161415625|
+-----+

```

```

In [4]: # if the given ip has a record in the following table
# the prediction for it's session length is the it's previous session's average
meanSessionIP = SessionDF.groupBy("ipaddress").agg(avg("SessionDuration").alias('AverageSessionDurationForIP'))
meanSessionIP.show(20)

```

```

+-----+-----+
|      ipaddress|AverageSessionDurationForIP|
+-----+-----+
| 27.62.30.188|33.0|
| 120.63.59.185|0.0|
| 115.69.247.81|30.0|
| 59.95.113.108|10.0|
|122.175.225.152|226.57692307692307|
| 14.139.60.13|0.0|
|117.232.164.217|1.5|
| 59.160.110.163|165.5|
| 101.62.250.135|27.0|
| 107.167.99.177|53.0|
| 121.246.85.180|4.0|
| 14.98.247.140|5.0|
| 1.23.208.26|10.833333333333334|
| 59.177.37.135|0.0|
|122.175.153.217|0.0|
| 223.176.174.84|2.0|
| 223.225.252.41|118.0|
|122.181.181.211|173.0|
| 103.15.63.34|46.666666666666664|
|123.238.121.215|3.0|
+-----+-----+

```

only showing top 20 rows

```
In [6]: # For Predicting the number of unique url visits for a given IP
        # again if we don't have the IP in our logs, the predicted value is the
        # average unique url visit by any IP
        SessionDF.groupBy("ipaddress", "url").count().distinct().groupBy().agg(avg("count")).show
```

```
+-----+
|      avg(count)|
+-----+
|1.3430700599135612|
+-----+
```

```
In [8]: # if we have the give ip in the records
        # we can find the average previous unique url visits of that IP in the following table
        SessionDF.groupBy("ipaddress", "url").count().distinct().groupBy("ipaddress").agg(avg("co
```

```
+-----+-----+
|      ipaddress|      avg(count)|
+-----+-----+
| 59.160.110.163|1.4285714285714286|
| 117.241.152.20|1.5588235294117647|
| 202.174.92.10|          1.0|
| 61.16.142.162|          1.0625|
| 117.205.39.248|          1.3125|
|117.203.181.144|          1.0|
|115.112.250.108|          1.0|
| 202.53.89.132|1.2272727272727273|
| 117.247.188.13|          1.0|
| 14.139.82.134|2.0555555555555554|
| 120.61.47.36|1.1428571428571428|
| 27.63.186.72|          1.0|
| 113.193.114.25|          1.0|
|123.136.182.137|          1.2|
| 27.34.244.251|1.1696428571428572|
| 124.125.22.218|          1.0|
| 117.207.97.173|          1.0|
| 61.0.225.164|1.3333333333333333|
|117.218.161.174|1.4210526315789473|
| 61.2.172.171|          1.0|
+-----+-----+
```

only showing top 20 rows

```
In [ ]:
```